

Fleetwright Technical Overview

A control plane for fleets of real browsers that books every job exactly once

Designed and built by Zahidul Islam

Evidence checked 6 October 2026

Fleetwright runs many logged-in Playwright browser sessions, watches a job feed, and books each matching job exactly once, even when workers crash, networks drop or two machines race for the same job. It is built as one shared core with separate worker roles: watchers that see new jobs, claimers that act on them, and a control process that dispatches, reconciles and measures. A planned crawler and a planned local AI browser agent reuse the same core.

The engine is tested end to end against a fictitious freight load board built in the same repository. The board has its own login with emailed one-time codes, a captcha, a live feed with ground-truth publish times, first-booker-wins booking and switches that inject failure on purpose. Fleetwright is never pointed at a real third-party site in this project, and this overview explains why that boundary is a legal requirement rather than a convenience.

Since 6 October 2026 the public site runs the whole system, not a prototype. Anyone can open the console and watch the engine claim loads from the mock board in real time; a shared demo login, shown on the sign-in page, can start a capped run and switch on failure modes, and everything resets automatically.

This overview is written for engineering reviewers, clients and collaborators. It separates what has been built and measured from what is still a design target. Credentials, private infrastructure addresses and personal filesystem paths are deliberately excluded. Source code and public documentation: github.com/ZahidulIslam2222/fleetwright. Live site: fleetwright.zahidul-islam.com.

1 The engineering problem

Teams that automate a web workflow with many accounts usually begin with independent scripts. That approach works for one or two browsers and fails at fifty. Two scripts book the same job, or neither does. A crash in the middle of a booking leaves nobody knowing whether it went through. Sessions expire, one-time codes arrive by email, and somebody has to sign in by hand. Nobody can say how quickly the fleet reacts, because nothing is measured.

The hard part is not driving a browser. It is coordination under failure: deciding which worker owns a job, rejecting a worker that woke up late, refusing to retry an action that may already have happened, and proving afterwards that every job was handled once. Fleetwright moves those decisions out of individual scripts and into one system with a database-enforced state

machine, encrypted session storage, a dispatcher with filters and schedules, a reconciler for uncertain outcomes, and per-stage latency measurement.

2 Scope and current status

The project is planned in ten phases, each closed by written acceptance criteria, automated gates, a real end-to-end check and an independent review. Phases one to five are complete: the design prototype and live site, the mock load board, the coordination core, the browser runtime and the live engine. Phase six, the live console, is shipped and running on the public site, with one item left: metrics and alerting. The crawler, the local AI agent and the large-scale proof runs are still ahead.

Area	Status
Mock load board with failure switches	Built and tested
Row-level security, claim state machine, fencing tokens, outbox	Built and tested
Worker runtime, session vault, one-time codes, failure evidence	Built and tested
Live engine: watchers, dispatcher, actors, reconciler, cells, latency	Built and tested
Console wired to the live engine	Built, tested and live
Console API, sign-in with password and TOTP, roles, audit log, live updates	Built, tested and live
Public read-only view and capped demo controls with automatic reset	Built, tested and live
Single-server production stack with rehearsed, hash-verified releases	Built and live
Metrics and alerting	Planned (remaining Phase 6 item)
Crawler, local AI agent, scale proof runs, cloud reference	Planned

3 Architecture and responsibilities

Postgres is the source of truth. It holds tenants, target-site accounts, sealed sessions, filters, schedules, jobs, claims with their full history, an outbox of messages to publish and an

append-only audit log. Every tenant table is protected by row-level security that is both enabled and forced, so even the table owner sees nothing without a tenant set.

Redis carries fast, short-lived traffic. Each cell has its own Redis with two streams per tenant: detected jobs and claims ready to act on. A core Redis holds control-plane keys such as console sessions. Losing Redis loses speed, never correctness, because every claim is already in Postgres and an outbox relay re-publishes anything left behind.

Worker processes run one Chromium each, with many isolated browser contexts called slots. A slot is either a watcher, which keeps the target's feed page open and listens to the JSON responses it fetches, or a claimer, which keeps a warm, logged-in page ready to book. A control process per cell runs the dispatcher, the outbox relay and the reconciler. Every service is stateless and can be killed at any time.

Component	Responsibility
Watcher slot	Turns each new load in the feed's responses into a detection with a seen time
Dispatcher	Drops fast-path duplicates, applies the tenant's filters and schedule windows, creates the claim and its outbox message in one transaction
Claimer slot	Leases the claim with a fencing token, acts only while the lease is current, records the answer
Reconciler	Applies lease timeouts and resolves uncertain claims by asking the target what the account booked
Outbox relay	Publishes anything the fast path missed, after checking the tenant still belongs to its cell
Console API	Serves the console's read models and owner actions under the caller's role and tenant
Live-update gateway	Holds one server-sent event stream per open console and tells it when something changed
Mock load board	The fictitious target with a ground-truth log used to audit every test

4 The claim state machine

A claim moves from queued to leased to acting, and then to confirmed, failed or unknown; an unknown claim ends as reconciled. Every transition is a single database statement, and each rule is enforced by the database rather than by worker code.

Exactly one claim exists per job, because a unique constraint on tenant and job key picks one winner however many watchers see the job. Leasing takes a fresh fencing token from a database sequence, and every later step names that token; a worker that paused during garbage collection or a network stall and woke up late is simply rejected. Lease expiry is judged only by the database clock, never by a worker clock.

The most important rule concerns crashes during the action. A claim whose lease expires while it is acting becomes unknown, and an unknown claim is never retried, because the booking may already have happened and a retry could book the job twice. The reconciler opens the account's sealed session, asks the target which loads the account holds, and records the truth. If the original worker later returns with a definite answer and a still-current token, that answer wins. A partial unique index also guarantees that one target account never holds two active leases.

5 Workers, sessions and one-time codes

Browser contexts are much cheaper than separate browsers, which is why each worker packs many of them into one Chromium. Accounts are handed to slots through a holder lease renewed on every heartbeat, so two workers never drive the same account. Sessions are stored as Playwright storage state, sealed per account, and refreshed before they expire. When a target asks for a one-time code, the worker reads it from email through an adapter, or waits for an operator to type it.

Secrets, cookies and one-time codes are never written to logs; a test scans the logs and published events to prove it. When a step fails, the worker records a browser trace and a screenshot with a retention limit, so failures can be investigated without guessing. Every read of a response body is bounded and every browser call has a deadline, because a hostile or broken page must not be able to freeze a slot.

6 Cells and moving tenants

A cell is one Redis plus the workers and control process that serve it. Each tenant belongs to exactly one cell, recorded in the database, so adding capacity means adding a cell and moving tenants, which is a data change rather than a code change. Redis keys carry the tenant as a hash tag, so a cell can later become a Redis Cluster without renaming anything.

Moving a tenant while the fleet is running is tested. The old cell's relay checks ownership inside its own transaction before publishing, the new cell picks up queued claims, and no job is booked twice. Claims already queued are handed over safely; detections in flight at the exact moment of the move may be dropped, which is recorded as a known limitation.

7 The live console and public demo

The console is a Next.js application exported as static files and served from the same origin as a versioned API, so there is no cross-origin traffic to secure. Every screen reads from the API through one strict contract: the server's response schemas are exported as an OpenAPI document, and a test fails if the server, the contract file or the console's role ranking drift apart.

There are four roles. Owners and operators can create and change filters and type one-time codes for target accounts; viewers can read everything; the demo role can only use the demo controls. Schedules are read-only in the console for now. A visitor who is not signed in gets a public, read-only view of the demo tenant with personal data masked, including the email addresses and network addresses in the audit log. Every change made through the console is written to the audit log with the actor, role and address range.

Live updates use server-sent events from a separate gateway. The stream carries only short hints that something changed, never the data itself; the console then refetches through the normal API, so a stream can never show a viewer data that their role should not see. A heartbeat keeps the connection open through proxies, and the console falls back to polling if the stream is unavailable.

The public demo is designed around abuse rather than politeness. The shared demo login and its current one-time code are shown on the sign-in page, so many visitors can sign in during the same 30 seconds. Its limits are enforced by the server, not by the screen: one run at a time, claimed atomically, with a maximum length; failure switches capped at safe values; and an automatic return to the calm preset fifteen minutes after the last change. The demo role cannot create or change filters or act on target accounts.

8 Production deployment

The public site runs on one shared server as a single container project: the web server, the console API, the live-update gateway, a control process, a browser worker, the mock board, Postgres, two Redis instances and a local mail catcher. Only the web server accepts connections, and only from the server's own reverse proxy behind a CDN. Every other service sits on an internal network with no route to the internet, runs with a read-only file system, no Linux capabilities and memory, CPU and process limits. The whole stack measured about 730 MB of memory in rehearsal.

Releases are immutable. Each one contains the static console, the rendered server configuration, the backend source at a single commit and a manifest of SHA-256 hashes. Before a release reaches the server it is started on a development PC as the exact production stack and checked end to end: pages and security headers, three demo visitors at once, capped controls, owner sign-in with TOTP, audit records, live updates and per-visitor limits. On the server, the live files are compared with their local copies before deploying and again afterwards; a deploy only counts as done when every hash matches.

The address of each visitor matters, because it drives rate limits, lockouts and the audit log. The proxy chain trusts only the hop it can see: the visitor's address reported by the CDN is believed only when the connection really comes from the CDN, and any address a client writes into its own headers is ignored. IPv6 visitors are counted per /64 block, the range one subscriber normally receives, so rotating through addresses does not escape the limits. A cap on total live streams sits below what the web server and the gateway can hold.

9 Security design

Security is designed in layers that are each tested. Tenant isolation relies on row-level security rather than on remembered filters in application code, and the test suite fails if forced security is removed from any table. Database roles are separated: an owner role for migrations, an application role, a worker role with column-level grants only, and a separate role for the mock board.

Target-site passwords and sessions use envelope encryption. Each record gets its own AES-256-GCM data key, wrapped by a master key, and the associated data binds the ciphertext to its tenant, account and purpose, so a sealed value copied into another row does not decrypt. In a cloud deployment the wrapping moves to a key management service without changing callers.

Console sign-in requires a password hashed with argon2id and a TOTP code for every account, with lockout per account and per network range and one-time use of each personal code. Password hashing runs outside the request loop with a fixed concurrency limit, and every password attempt from one network range counts towards a window, because the demo password is public and a correct password would otherwise be free work for an attacker. Sessions use secure, HTTP-only cookies and every change needs a CSRF token. The console ships with a content security policy and related headers. Every file edit is scanned for secrets and Python security issues, every commit also runs semgrep, and the same scanners run in GitHub Actions. Findings are fixed, never suppressed. All SQL is static and parameterised.

Threat	Response	Status
Tenant data crossing tenants	Forced row-level security, tested and mutation-checked	In place
Stale worker acting on a lost claim	Fencing tokens and database-clock leases	In place
Database dump exposing sessions	Envelope encryption with row-bound associated data	In place
Secrets in logs or source control	Log-leak test; secret scanning on edit, commit and push	In place
Operator account takeover	Argon2id, mandatory TOTP, lockout, replay protection, audit log	In place, live
Public demo login used to exhaust the server	Bounded hashing, per-range attempt window, per-visitor limits	In place, live
Forged client address to dodge limits or poison the audit log	Trust only the visible hop; CDN header only from the CDN	In place, tested
Address rotation to bypass limits	IPv6 counted per /64; total stream cap	In place, tested
Server-side request forgery from crawl targets	Egress guard after DNS resolution	Planned
Prompt injection into the AI agent	Human approval before publish, send, purchase or delete	Planned

10 Law and acceptable use

The legal boundary shaped the project as much as the technology. In the United States, courts have read the Computer Fraud and Abuse Act narrowly for public pages since *Van Buren* and *hiQ v. LinkedIn*, and the fight has moved to contract law. *Meta v. Bright Data* held that logged-out scraping of public data did not breach Meta's terms, but terms do bind logged-in use, which is exactly what a booking fleet does. Major freight load boards forbid automation outright:

DAT's terms, last updated in July 2026, prohibit bots, AI agents and browser extensions used to access the product without written approval and allow immediate termination.

Fleetwright therefore ships with its own fictitious board, contains no code that defeats third-party captchas or access controls, and requires a written onboarding gate before any real target: the client must show that the platform permits automation or that it owns the platform. The live public demo runs only against that built-in board, on an internal network that cannot reach the internet. In the European Union the design follows GDPR minimisation, retention and data-subject rights with Fleetwright acting as processor, uses strictly necessary cookies only, plans AI Act transparency labelling for the agent, and prepares for Cyber Resilience Act vulnerability reporting with a published security policy. The full register is in the repository's legal documentation and is explicitly not legal advice.

11 Scaling to one million sessions and 99.9 percent availability

Two different scales matter. The control plane, meaning the API, the console and live updates, is designed for one million concurrent connected sessions across many tenants. The worker plane is designed to grow from tens to more than ten thousand browser contexts by adding cells. A million automated browsers would need roughly one hundred terabytes of memory, which is a different problem that the design does not claim to solve.

The availability target for the cloud reference architecture is 99.9 percent per month, about 43 minutes of downtime, which is stricter than the 99 percent originally requested (about 7.2 hours). The public demo runs on a single shared server with no redundancy and no backup job yet, so it claims no availability figure at all. Correctness is not treated as a service level: a single duplicate booking is an incident, never part of an error budget.

Stage	Shape	Capacity	Availability
Today	One shared server running the live demo; development PC for load tests	4 live slots on the demo; measured up to 32 contexts on a PC	No claim
Single production server	Dedicated server with backups and monitoring	About 75 contexts per 8 GB node (estimate)	About 99 percent with good operations (estimate)
Cloud, one region, multi-AZ	Managed Postgres and Redis, autoscaled	Hundreds of contexts, thousands of console sessions	99.9 percent target

Stage	Shape	Capacity	Availability
	containers, load balancer and firewall		
Many cells	Tenants spread across cells, gateway fleet	More than 10,000 contexts (target)	99.9 percent per cell
One million sessions	Gateway fleet, CDN for the console, read replicas, partitioned history tables	1,000,000 connected sessions (target)	99.9 percent; 99.95 needs multi-region

The seams that make this possible are already in the code and tested: tenant identifiers leading every key and index, stateless services, cell routing, hash-tagged Redis keys, an outbox with idempotent consumers, bounded queries and one typed configuration boundary. Kubernetes, Kafka, multi-region deployment and database sharding are deliberately not built yet; each has a written trigger that would justify it. The gateway's capacity per process has not been load-tested yet.

12 Verification and measured results

Correctness is checked against the board's own ground-truth log, so the system is never marking its own homework. Ninety-nine automated tests pass against real Postgres, Redis, a mail server, a real Chromium and the real mock board, alongside type checking, linting and security scanning on every commit.

Measurement	Result
Duplicate bookings with 200 concurrent competitors and 10,000 jobs	0; 9,897 claims confirmed equal 9,897 bookings on the board
Crashes injected during the load test	184 after leasing and 213 during the action, all re-queued or reconciled
Chaos run with lost acknowledgements, errors, slow responses and rival bots	Zero duplicates; every lost answer reconciled
Detection to confirmed booking	Median about 365 ms, 95th percentile about 540 ms

Measurement	Result
Memory per browser context	About 70 to 103 MB; 100 MB used for planning
CPU at 32 contexts	About 1.8 cores
Tenant isolation	Tenant A cannot read or write tenant B rows, even as the table owner
Live site end-to-end check through the CDN	0 failures: three demo visitors at once, capped controls, owner sign-in with TOTP, audit, live updates
Live update reaching a client through the CDN	First engine update 5.8 seconds after connecting (one live measurement, with the demo feed at its idle rate)
Per-visitor limits on the live site	A 160-request burst from one visitor: 71 answered, 89 refused with 429
Client address tests in rehearsal	Forged headers ignored in every case; IPv6 visitors grouped per /64
Release integrity	Every deployed file matches its recorded hash before and after the deploy
Console data on screen, repeat visit	About 1.4 seconds, measured from Bangladesh

Several checks were broken on purpose to prove they could fail: removing forced row-level security from one table fails the isolation suite, writing a secret into a log line fails the leak test, and undoing the shared demo login fix makes the three-visitor test fail. The large measurements come from one developer PC against the mock board; the 50 and 70 context one-hour runs are still to be done.

Each phase also had a review by a reviewer that did not write the code. For the live console, the first review said not to ship: the shared demo login let in only one visitor per 30 seconds, password hashing could stall the API, and there were no per-visitor limits. All three were fixed with tests. The second review found that IPv6 visitors could avoid the limits; that was fixed and tested before the deploy.

13 Failures and the engineering lessons

The reconciler's sweep blocked workers under load. In the 10,000 job test, worker lease updates hit the five-second statement timeout because the sweep moved many claims in one statement and then wrote history rows one by one inside the same transaction, holding locks for seconds. Each sweep rule became a single statement that moves rows and writes history and outbox entries together, and the reconciler stopped holding locks while it called the target. Lesson: never hold a lock across a loop of round trips, and never across a network call.

A green concurrency test proved almost nothing. The first passing run booked only 77 of 10,000 jobs: leases were so short that jobs aged out before anyone acted, and "zero duplicates" was trivially true. Production-like timings and coverage assertions were added, requiring at least 95 percent of jobs booked and at least one crash during the action reconciled. Lesson: a passing test must also prove it exercised the dangerous path.

A browser call with no timeout froze a worker slot. About one run in three, after the target revoked a session, the slot never signed in again and logged nothing. A dump of every suspended task pointed at a page close call, which has no timeout, waiting on a page that was redirecting itself. Sign-in now uses a new page and closes the old one in a bounded background task. Lesson: every call into a browser needs a deadline, especially the ones without a timeout parameter.

An old cell could publish work after its tenant had moved away. The relay's cached tenant list was a few hundred milliseconds stale, so it published one message to the wrong Redis. The relay now checks ownership inside its own transaction with a share lock that a move must wait for. Lesson: a cached ownership list is a hint; the check that guards a side effect belongs in the same transaction.

Click mode froze claimers. A response body read never returned, and the button lookup ran after the claim entered the acting state, so a slow lookup plus the click could outlive the lease. Body reads are now bounded, harmless steps run before acting, and the fleet refuses to start unless the lease outlasts the action timeout; that check immediately caught a wrong setting in the test kit. Lesson: do the slow, harmless steps before taking a lock that has a deadline, and make the deadline relationship a startup check.

The first capacity number was wrong by a factor of two. Summing resident memory across Chromium processes counted shared memory many times, the measurement included the test's own board server, and CPU always read zero because fresh process objects were created for every sample. Unique memory per process and persistent process objects gave 70 to 103 MB per context and about 1.8 cores at 32 contexts, and the same CPU bug was fixed in the worker heartbeat. Lesson: before planning capacity on a number, check what the meter actually counts.

The public demo login let in one visitor per 30 seconds. One-time codes are single use per account, which is right for a personal account and wrong for a shared one whose code is published: the second visitor in the same 30 seconds was rejected, and each retry counted towards a lockout. My own browser check signed in once, so it passed. The shared demo account now skips the per-account replay check while every personal account keeps it, and a test signs in three visitors with the same code. Lesson: a security control that is right for one user can be a denial of service for a shared account, so test the shared path with more than one user.

My own design for reading the visitor's address could be forged. The first proxy configuration skipped every forwarded hop inside the CDN's published address ranges to find the visitor, but code running on the CDN's own platform sends requests from those same ranges, so anyone could relay through it and choose the address that rate limits and the audit log would believe. The proxy now reads only the hop it can see and accepts the CDN's visitor header only when that hop is the CDN. Lesson: trusted address ranges are only as trusted as everyone else who can send from them.

A security test passed because it could not see. A check that a backend container could not forge its address reported success, but the audit value it read was "hidden": its owner sign-in had silently reused a one-time code an earlier script had spent, so it read the log as an anonymous visitor with addresses masked. The scripts now require the sign-in to succeed and wait for a fresh code, and the rerun showed the real range. Lesson: assert a security test's preconditions, or a negative check passes for the wrong reason.

The production smoke test was quietly steering the live engine. It proved that an owner can create a filter by creating one, and the API has no delete, so every run would have added another active rule to the engine visitors were watching. The test now reuses one filter and leaves it disabled. Lesson: a smoke test against production is a production user and has to clean up like one.

14 Roadmap and known limitations

The remaining Phase 6 item is metrics and alerting. After that come a crawler with politeness, robots and opt-out handling and server-side request forgery protection; a lightweight local AI browser agent with human approval for risky actions; and a hardening phase with an OWASP ASVS Level 2 checklist, dynamic scanning, software bills of materials and signed images, the 50 and 70 context one-hour fleet runs, a gateway load test, a restore drill and a cloud reference architecture in infrastructure as code.

Known limitations are recorded rather than hidden. Blocked responses mark a claim as failed instead of re-queuing it to another account. Per-account rate pacing and routing by account group are configured but not yet enforced by the dispatcher. Detections in flight during a tenant move may be dropped. Re-planning after a tenant change restarts every slot in that worker. The

public site runs on one server with no backup job yet; the demo data is re-created by a seed job. There are no metrics or alerts yet. The audit entry for a console action is written in its own transaction just after the action. A revoked session keeps an already-open live-update stream until it reconnects. The content security policy still allows inline scripts and styles. The extra load from many open consoles refetching on each update has not been load-tested, and the tests share the local development databases.

15 Source basis

This edition draws on the Fleetwright source code, its automated test suite and test output, the project's measurement runs, the release rehearsal and live-site checks, its architecture decision records and public documentation, and the legal sources cited in the repository's compliance register, checked on 5 and 6 October 2026. Private planning notes, credentials and infrastructure details remain in private project records.

The strongest claims are supported by executable tests against real infrastructure and a ground-truth log kept by the target itself. Design targets for scale and availability are labelled as targets throughout and are not presented as achieved capacity.